

Dimension Reduction

David Hofmeyr

Dept. Statistics and Actuarial Science,
Stellenbosch University, South Africa

6-8 December 2021



forward together · saam vorentoe · masiye pbambili

Projection Pursuit

Independent Components Analysis

Projection Pursuit

- ▶ (Usually) linear dimension reduction or feature extraction
- ▶ Often unsupervised: data exploration
- ▶ Look for “interesting projections” of a given set of data
- ▶ What is interesting?
- ▶ What's your objective?
 - ▶ Clustering? clusters are interesting
 - ▶ Outlier detection? outliers are interesting
 - ▶ Regression? Derived covariates with a close predictive relationship to the response are interesting
 - ▶ Classification? simple boundaries between classes are interesting
 - ▶ No idea? we can help you too
 - ▶ High information projections

Projection Pursuit

- Broadly speaking, from a mathematical point of view:

$$\max_{\mathbf{V} \in \mathbb{R}^{p \times p'}} \Phi(\mathbf{V}|\mathbf{X}),$$

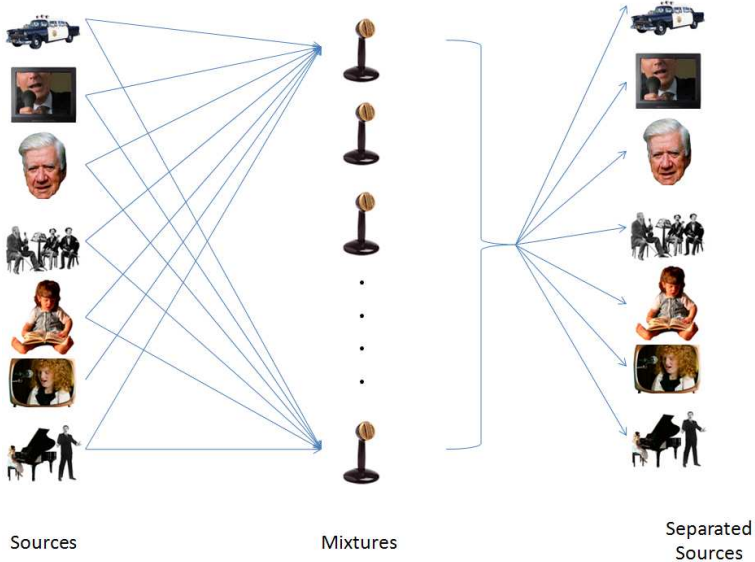
where Φ measures the “interestingness” of the projected data \mathbf{XV} .

- Ultimately we are interested in being able to compute some measure of interestingness (whatever that means) of the distribution of the projected data
- The more detail we estimate in this distribution, the more complex our measures can be
 - Second moment: PCA
 - Third moment ...
 - Fourth moment:
 - minimise \rightarrow clustering
 - maximise \rightarrow outliers
 - etc.

Projection Pursuit

- ▶ at the extreme: exact estimates of arbitrary functionals of $F_{\mathbf{X}\mathbf{v}}$ or $F_{Y|\mathbf{X}\mathbf{v}}$
- ▶ I live at this extreme
- ▶ (Very fortunately) frequently the objective can be decomposed over the different *components*, $\mathbf{X}\mathbf{v}_1, \dots, \mathbf{X}\mathbf{v}_{p'}$
 - ▶ or we force it to decompose since the optimisation problem is considerably easier
- ▶ We still are faced with a computational problem:
 - ▶ Estimate arbitrary functional from the distribution of univariate (projected) data $\mathbf{X}\mathbf{v}$
 - ▶ Compute the gradient of this functional w.r.t. \mathbf{v}
 - ▶ Repeatedly, over multiple *projection vectors*, $\mathbf{v}^{(0)}, \mathbf{v}^{(1)}, \mathbf{v}^{(2)}, \dots$ during optimisation

Independent Component Analysis: The Cocktail Party Problem



Independent Component Analysis: The Model

- ▶ Identify *source signals* which are observed only indirectly after “mixing”
- ▶ Assume $\mathbf{S} \in \mathbb{R}^{n \times k}$ represents realisations of k INDEPENDENT source signals
- ▶ Assume observations are given by $\mathbf{X} = \mathbf{SM}$
- ▶ Task is to estimate *unmixing* matrix, \mathbf{U} , s.t., $\mathbf{XU} = \mathbf{CPS}$
 - ▶ \mathbf{C}, \mathbf{P} are inestimable scaling and permutation matrices
 - ▶ We estimate the sources up to arbitrary re-ordering and scalar multiplication
- ▶ How could we possibly start to address this?
 - ▶ We assume nothing about the sources (except independence)
- ▶ We could use projection pursuit to “maximise independence”
 - ▶ what do we even mean by this?
 - ▶ measuring independence is not usually straight-forward (and is computationally demanding)

Independent Component Analysis: The Method

- ▶ Some useful observations massively simplify the problem formulation
 - ▶ Orthogonality (zero covariance) is necessary for independence
 - ▶ scale doesn't affect independence: $X \perp Y \Rightarrow aX \perp Y \forall a$
- ▶ We can measure the independence in the components of \mathbf{XV} via their mutual information

$$\begin{aligned} MI(\mathbf{XV}) &= KL \left(f_{\mathbf{XV}} \parallel \prod_i f_{\mathbf{Xv}_i} \right) \\ &= E_{Z \sim F_{\mathbf{XV}}} [\log(f_{\mathbf{XV}}(Z))] - E_{Z \sim F_{\mathbf{XV}}} \left[\log \left(\prod_i f_{\mathbf{Xv}_i}(Z_i) \right) \right] \\ &= E_{Z \sim F_{\mathbf{X}}} [\log(f_{\mathbf{X}}(Z))] + \log(|\det(\mathbf{V})|) \\ &\quad - \sum_i E_{Z \sim F_{\mathbf{Xv}_i}} [\log(f_{\mathbf{Xv}_i}(Z))] \end{aligned}$$

for non-singular \mathbf{V} .

Independent Component Analysis

- ▶ But $E_{Z \sim F_{\mathbf{X}}}[\log(f_{\mathbf{X}}(Z))]$ is constant, and we don't care about scale (we can't discriminate based on scale), so can force $\det(\mathbf{V})$ to be constant
- ▶ We therefore want to minimise

$$- \sum_i E_{Z \sim F_{\mathbf{X}_{\mathbf{v}_i}}}[\log(f_{\mathbf{X}_{\mathbf{v}_i}}(Z))]$$

- ▶ We don't know $f_{\mathbf{X}_{\mathbf{v}_i}}$:
 - ▶ we assume we have a sample from $F_{\mathbf{X}}$ for estimation
- ▶ minimise the sample estimate

$$- \sum_i \frac{1}{n} \sum_{j=1}^n \log(\hat{f}_{\mathbf{X}_{\mathbf{v}_i}}(\mathbf{x}_j^\top \mathbf{v}_i))$$

- ▶ Notice that this is also the maximum (pseudo) likelihood solution, under the assumption of independence

A Quick Aside on Information

- ▶ A random variable carrying a lot of information is one which “has a tendency to arise in high density regions”
 - ▶ Why is this “information” / “informative”
 - ▶ Has a tendency to “show you” where independent copies of itself are likely to “land”
- ▶ ALL random variables have “tendency to arise in high density regions”
- ▶ Some to a greater extent: For non-decreasing $m : \mathbb{R} \rightarrow \mathbb{R}$,

$$E_X[m(f_X(X))] \text{ large} \Rightarrow \text{Information}(X) \text{ large.}$$

- ▶ ICA \equiv maximise the Shannon information of the projections

Independent Component Analysis: The Computation

- ▶ Computational problem: Estimating $f_{\mathbf{x}\mathbf{v}}$, and evaluating it at all $\mathbf{x}_j^\top \mathbf{v}$ is expensive if approached naïvely
 - ▶ Approximations from “negentropy” (measure of departure from Gaussianity)
 - ▶ ... or... or
 - ▶ or suck it up and get better at computation
- ▶ To be completely general, we must be non-parametric
 - ▶ I like kernels, they’re more intuitive (to me) and amenable to differentiation
- ▶ The kernel estimate of a density, f_X , using a sample from its distribution, $\{x_1, \dots, x_n\}$, is given by

$$\hat{f}(x) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - x_i}{h}\right)$$

Fast Kernel Smoothing

- We'll be slightly more general, and look at evaluating

$$\hat{f}(x) = \sum_{i=1}^n K\left(\frac{x - x_i}{h}\right) \omega_i,$$

for arbitrary $\{\omega_1, \dots, \omega_n\}$

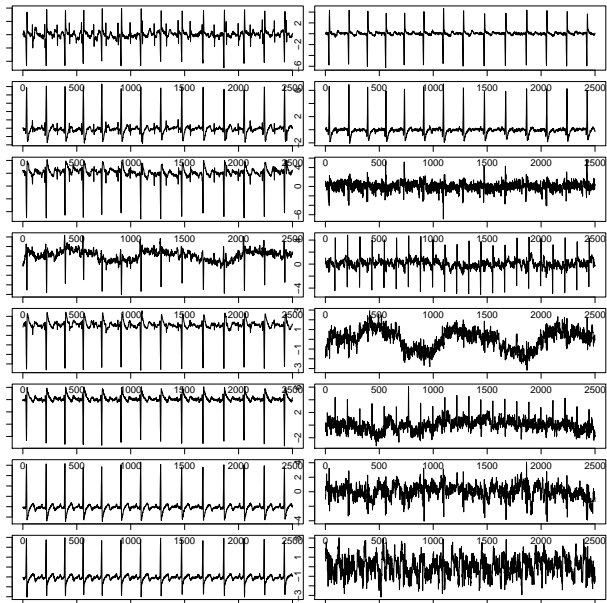
- To evaluate directly, costs $\mathcal{O}(nm)$ to evaluate at m points
- But, if the kernels look like this:

$$K(x) = \sum_{i=0}^{\alpha} \beta_i |x|^i \exp(-|x|)$$

we can do much better ($\mathcal{O}(n \log(n))$)

- Fast computation boils down to:
 1. trivial factorisation of $\exp(a + b)$
 2. binomial expansion of $(a + b)^i$


A Few ICA Successes: Foetal ECG



A Few ICA Successes: Reflection removal



ICA: What's it actually doing? (vs PCA)



Other Examples: Projection Pursuit for Clustering


- Maximise over \mathbf{v} ,

$$\Phi(\mathbf{v}|\mathbf{X}) = \max_b Q(\mathbf{v}, b|\mathbf{X}),$$


where Q measures the quality of clustering \mathbf{X} using hyperplane $H(\mathbf{v}, b) = \{\mathbf{x} \in \mathbb{R}^p | \mathbf{x}'\mathbf{v} = b\}$.

- Examples:
 - Minimise surface integral of density on $H(\mathbf{v}, b)$
 - Minimise normalised cut across $H(\mathbf{v}, b)$
 - Maximise variance ratio of clusters (similar to LDA objective)


Other Examples: Projection Pursuit for Clustering (in action)




Other Examples: Projection Pursuit for Clustering (in action)




Other Examples: Projection Pursuit for Clustering (in action)




Other Examples: Projection Pursuit for Clustering (in action)




Other Examples: Projection Pursuit for Clustering (in action)




Other Examples: Projection Pursuit for Clustering (in action)




Other Examples: Projection Pursuit for Clustering (in action)




Other Examples: Projection Pursuit for Clustering (in action)




Other Examples: Projection Pursuit for Clustering (in action)




Other Examples: Projection Pursuit for Clustering (in action)



Other Examples: Projection Pursuit for Clustering (in action)



Other Examples: Projection Pursuit for Clustering (in action)



Other Examples: Projection Pursuit Regression

- ▶ Additive model with “ridge” functions
- ▶ GAM with derived features
- ▶ single layer NN with non-parametric activation function
- ▶ Estimate $f(\mathbf{x}) = E[Y|\mathbf{X} = \mathbf{x}]$ by minimising

$$\sum_{i=1}^n \mathcal{L} \left(y_i, \mu + \sum_{j=1}^k g_j(\mathbf{x}_i^\top \mathbf{v}_j) \right)$$

Other Examples: Projection Pursuit Regression

- Each g_j is fit non-parametrically (e.g., with kernels)
- For notational brevity: $S_j(x|\mathbf{b}) = \sum_{i=1}^n K\left(\frac{\mathbf{x}_i^\top \mathbf{v}_j - x}{h}\right) b_i$
- Local constant (Nadaraya-Watson):

$$\hat{g}_j(x) = \frac{\sum_{i=1}^n K\left(\frac{\mathbf{x}_i^\top \mathbf{v}_j - x}{h}\right) y_i}{\sum_{i=1}^n K\left(\frac{\mathbf{x}_i^\top \mathbf{v}_j - x}{h}\right)} = \frac{S_j(x|\mathbf{y})}{S_j(x|\mathbf{1})}$$

- Local linear:

$$\begin{aligned}\hat{g}_j(x) = & \frac{S_j(x|(\mathbf{X}\mathbf{v}_j)^2)S_j(x|\mathbf{y}) - S_j(x|\mathbf{X}\mathbf{v}_j)S_j(x|\mathbf{X}\mathbf{v}_j * \mathbf{y})}{S_j(x|(\mathbf{X}\mathbf{v}_j)^2)S_j(x|\mathbf{1}) - S_j(x|\mathbf{X}\mathbf{v}_j)^2} \\ & + \frac{S_j(x|\mathbf{1})S_j(x|\mathbf{X}\mathbf{v}_j * \mathbf{y}) - S_j(x|\mathbf{X}\mathbf{v}_j)S_j(x|\mathbf{y})}{S_j(x|(\mathbf{X}\mathbf{v}_j)^2)S_j(x|\mathbf{1}) - S_j(x|\mathbf{X}\mathbf{v}_j)^2} x\end{aligned}$$

Other Examples: Optimal Projections for NB

- ▶ Naïve Bayes has enjoyed a lot of success, despite its very simplistic model: Classify based on standard

$$P(Y = k | \mathbf{X} = \mathbf{x}) = \frac{f_{\mathbf{x}|Y=k}(\mathbf{x})P(Y = k)}{\sum_{j=1}^K f_{\mathbf{x}|Y=j}(\mathbf{x})P(Y = j)}$$

but treat $f_{\mathbf{x}|Y=k}$ as though

- ▶ Treat covariates as being completely independent, conditional on the class label
- ▶ In the Gaussian case, like diagonal discriminant analysis
- ▶ By treating these as independent, we can fit more flexible models (e.g., non-parametric ones) easily
- ▶ Can we find a projection \mathbf{V} , so that NB applied to \mathbf{XV} fits well?
 - ▶ Like changing the basis for factorisation of the densities

Other Examples: Optimal Projections for NB

- ▶ An earlier approach applied ICA to each class separately (CCICA) to better approximate the independence assumptions: Hopefully less bias in the estimate of each $f_{\mathbf{x}|Y=k}$
 - ▶ May not aid in better discrimination
- ▶ Maximise the classification likelihood (multinomial likelihood, conditional on fixed \mathbf{X}), with a single projection matrix

$$\prod_{i=1}^n \frac{\hat{f}_{\mathbf{xv}|Y=y_i}(\mathbf{x}_i^\top \mathbf{v}) \hat{P}(Y = y_i)}{\sum_{j=1}^K \hat{f}_{\mathbf{xv}|Y=j}(\mathbf{x}_i^\top \mathbf{v}) \hat{P}(Y = j)}$$


Other Examples: Optimal Projections for NB

- (Obviously) we log this, and then do some simple algebra, and find the objective is given by

$$\sum_{k=1}^K \left(\log \left(\frac{n_k}{n} \right) + \sum_{j: y_j = k} \log \left(\hat{f}_{\mathbf{x}\mathbf{v}|Y=k}(\mathbf{x}_j^\top \mathbf{v}) \right) \right) - \sum_{j=1}^n \log \left(\sum_{k=1}^K \frac{n_k}{n} \hat{f}_{\mathbf{x}\mathbf{v}|Y=k}(\mathbf{x}_j^\top \mathbf{v}) \right)$$

- The first term (after factorising the densities) is a weighted sum of the ICA objectives for the classes, and the second term penalises for non-discrimination
- This can similarly be written using combinations of kernel weighted sums

Some Plots (vs Others)



Playing with ICA in R

- ▶ We will implement our own ICA method in R using kernel estimates for the densities
- ▶ The package FKSUM provides implementations of the log-linear time methods
- ▶ We will use a similar approach to that used in PCA based on projection pursuit, where projection of the argument onto the unit ball forms part of the evaluation: If $\Phi(\mathbf{v}|\mathbf{X}) = \phi(\mathbf{X}\vec{\mathbf{v}})$ is the objective function, then

$$\nabla_{\mathbf{v}}\Phi(\mathbf{v}|\mathbf{X}) = \frac{1}{\|\mathbf{v}\|} (I - \vec{\mathbf{v}}\vec{\mathbf{v}}')\mathbf{X}'\nabla_{\mathbf{p}}\phi(\mathbf{p})\big|_{\mathbf{p}=\mathbf{X}\vec{\mathbf{v}}}$$

► We have

$$\begin{aligned}\phi(\mathbf{p}) &= -\frac{1}{n} \sum_{i=1}^n \log(\hat{f}_{\mathbf{p}}(p_i)) = -\frac{1}{n} \sum_{i=1}^n \log \left(\frac{1}{nh} \sum_{j=1}^n K \left(\frac{p_i - p_j}{h} \right) \right) \\ \Rightarrow \frac{\partial}{\partial p_k} \phi(\mathbf{p}) &= \frac{1}{n^2 h^2} \sum_{i=1}^n K' \left(\frac{p_i - p_k}{h} \right) \left(\frac{1}{\hat{f}_{\mathbf{p}}(p_i)} + \frac{1}{\hat{f}_{\mathbf{p}}(p_k)} \right)\end{aligned}$$